

```

import os, sys
import pandas as pd
import numpy as np
from pathlib import Path

fp = Path("weather_forecast_data.csv")
print("File exists:", fp.exists(), "size:", fp.stat().st_size if
fp.exists() else None)

df = pd.read_csv(fp)
print("Loaded dataframe with shape:", df.shape)
df.head(10)

```

File exists: True size: 251223
Loaded dataframe with shape: (2500, 6)

	Temperature	Humidity	Wind_Speed	Cloud_Cover	Pressure	
Rain						
0	23.720338	89.592641	7.335604	50.501694	1032.378759	
rain						
1	27.879734	46.489704	5.952484	4.990053	992.614190	no
rain						
2	25.069084	83.072843	1.371992	14.855784	1007.231620	no
rain						
3	23.622080	74.367758	7.050551	67.255282	982.632013	
rain						
4	20.591370	96.858822	4.643921	47.676444	980.825142	no
rain						
5	26.147353	48.217260	15.258547	59.766279	1049.738751	no
rain						
6	20.939680	40.799444	2.232566	45.827508	1014.173766	no
rain						
7	32.294325	51.848471	2.873621	92.551497	1006.041733	no
rain						
8	34.091569	48.057114	5.570206	82.524873	993.732047	no
rain						
9	19.586038	82.978293	5.760537	98.014450	1036.503457	
rain						

```

# Full analytics pipeline on the loaded dataframe
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
import joblib

# show basic info
info = {}

```

```
info['dtypes'] = df.dtypes.to_dict()
info['missing'] = df.isnull().sum().to_dict()
info['describe'] = df.describe().T
```

```
print("Column dtypes:\n", info['dtypes'])
print("\nMissing values:\n", info['missing'])
print("\nSummary statistics:\n", info['describe'].head(10))
```

Column dtypes:

```
{'Temperature': dtype('float64'), 'Humidity': dtype('float64'),
'Wind_Speed': dtype('float64'), 'Cloud_Cover': dtype('float64'),
'Pressure': dtype('float64'), 'Rain': dtype('O')}
```

Missing values:

```
{'Temperature': 0, 'Humidity': 0, 'Wind_Speed': 0, 'Cloud_Cover': 0,
'Pressure': 0, 'Rain': 0}
```

Summary statistics:

	count	mean	std	min	25%
Temperature	2500.0	22.581725	7.326996	10.001842	16.359398
Humidity	2500.0	64.347094	19.954739	30.005071	47.339815
Wind_Speed	2500.0	9.906255	5.780316	0.009819	4.761909
Cloud_Cover	2500.0	49.658104	29.123104	0.015038	23.900016
Pressure	2500.0	1014.312336	20.196433	980.014486	996.938630

	50%	75%	max
Temperature	22.536448	28.976476	34.995214
Humidity	63.920797	81.561021	99.997481
Wind_Speed	9.908572	14.948408	19.999132
Cloud_Cover	49.488284	75.324140	99.997795
Pressure	1013.433035	1031.735067	1049.985593

Convert target to binary

```
df_clean = df.copy()
if 'Rain' in df_clean.columns:
    le = LabelEncoder()
    df_clean['Rain_bin'] = le.fit_transform(df_clean['Rain'])
    print("\nLabel classes (Rain):", list(le.classes_))
else:
    raise ValueError("No 'Rain' column found to serve as target.")
```

Label classes (Rain): ['no rain', 'rain']

Check for non-numeric columns besides target

```
non_numeric =
df_clean.select_dtypes(include=['object']).columns.tolist()
print("\nNon-numeric columns:", non_numeric)
```

If any numeric-like objects, try to coerce

```

for col in non_numeric:
    if col != 'Rain':
        try:
            df_clean[col] = pd.to_numeric(df_clean[col],
errors='coerce')
        except:
            pass

```

Non-numeric columns: ['Rain']

```

#Drop rows with missing values (report how many)
before = len(df_clean)
df_clean = df_clean.dropna()
after = len(df_clean)
print(f"\nDropped {before-after} rows with missing values. Remaining
rows: {after}")

```

```

# Correlation matrix
corr = df_clean.select_dtypes(include=[np.number]).corr()
print("\nNumeric correlation:\n", corr)

```

Dropped 0 rows with missing values. Remaining rows: 2500

Numeric correlation:

	Temperature	Humidity	Wind_Speed	Cloud_Cover	Pressure
Temperature	1.000000	-0.014119	-0.004957	0.006632	-0.007489
Humidity	-0.014119	1.000000	0.003816	0.007244	-0.032089
Wind_Speed	-0.004957	0.003816	1.000000	-0.000405	0.018772
Cloud_Cover	0.006632	0.007244	-0.000405	1.000000	0.023100
Pressure	-0.007489	-0.032089	0.018772	0.023100	1.000000
Rain_bin	-0.265882	0.382464	-0.001389	0.326216	0.008273

	Rain_bin
Temperature	-0.265882
Humidity	0.382464
Wind_Speed	-0.001389
Cloud_Cover	0.326216
Pressure	0.008273
Rain_bin	1.000000

```

# Show top 10 rows to user

```

```
try:
    from ace_tools import display_dataframe_to
    display_dataframe_to("Sample of cleaned data", df_clean.head(50))
except Exception:
    # fallback: print head
    print("\nSample of cleaned data:\n", df_clean.head(10))
```

Sample of cleaned data:

	Temperature	Humidity	Wind_Speed	Cloud_Cover	Pressure	
Rain \						
0	23.720338	89.592641	7.335604	50.501694	1032.378759	
rain						
1	27.879734	46.489704	5.952484	4.990053	992.614190	no
rain						
2	25.069084	83.072843	1.371992	14.855784	1007.231620	no
rain						
3	23.622080	74.367758	7.050551	67.255282	982.632013	
rain						
4	20.591370	96.858822	4.643921	47.676444	980.825142	no
rain						
5	26.147353	48.217260	15.258547	59.766279	1049.738751	no
rain						
6	20.939680	40.799444	2.232566	45.827508	1014.173766	no
rain						
7	32.294325	51.848471	2.873621	92.551497	1006.041733	no
rain						
8	34.091569	48.057114	5.570206	82.524873	993.732047	no
rain						
9	19.586038	82.978293	5.760537	98.014450	1036.503457	
rain						

	Rain_bin
0	1
1	0
2	0
3	1
4	0
5	0
6	0
7	0
8	0
9	1

Plots: histograms for numeric columns

num_cols =

df_clean.select_dtypes(include=[np.number]).columns.tolist()

exclude target

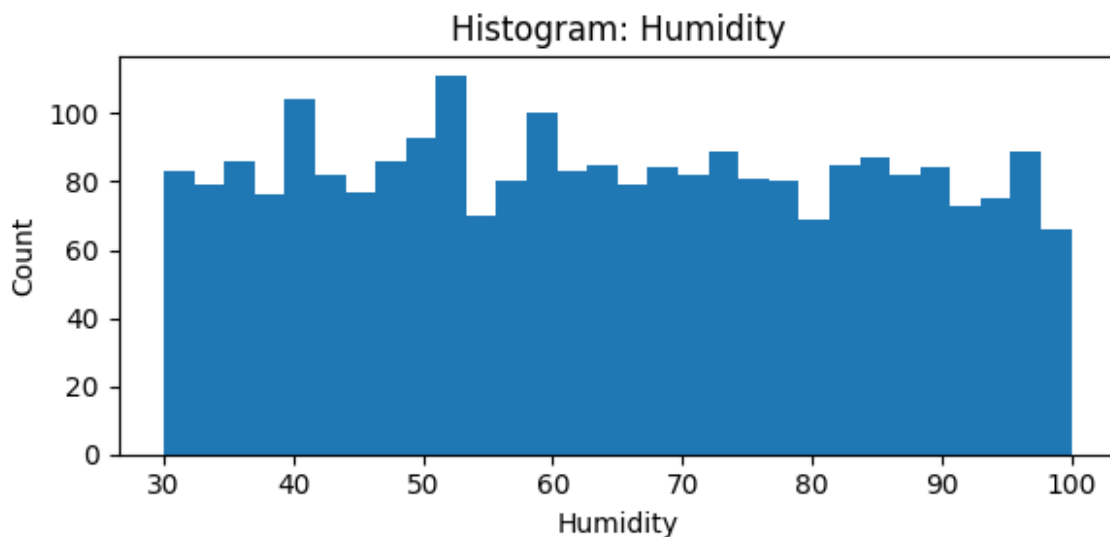
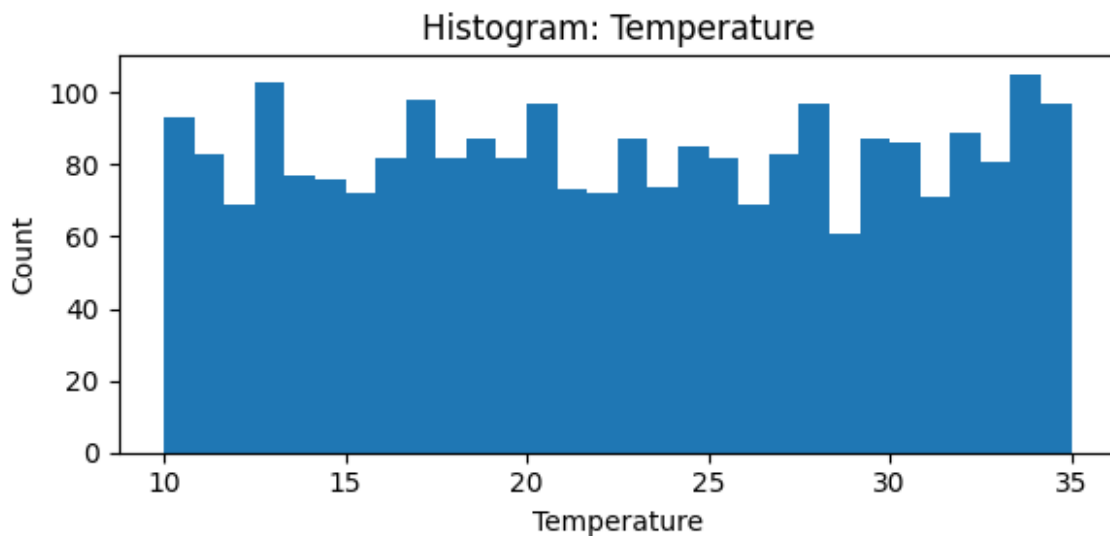
if 'Rain_bin' in num_cols:

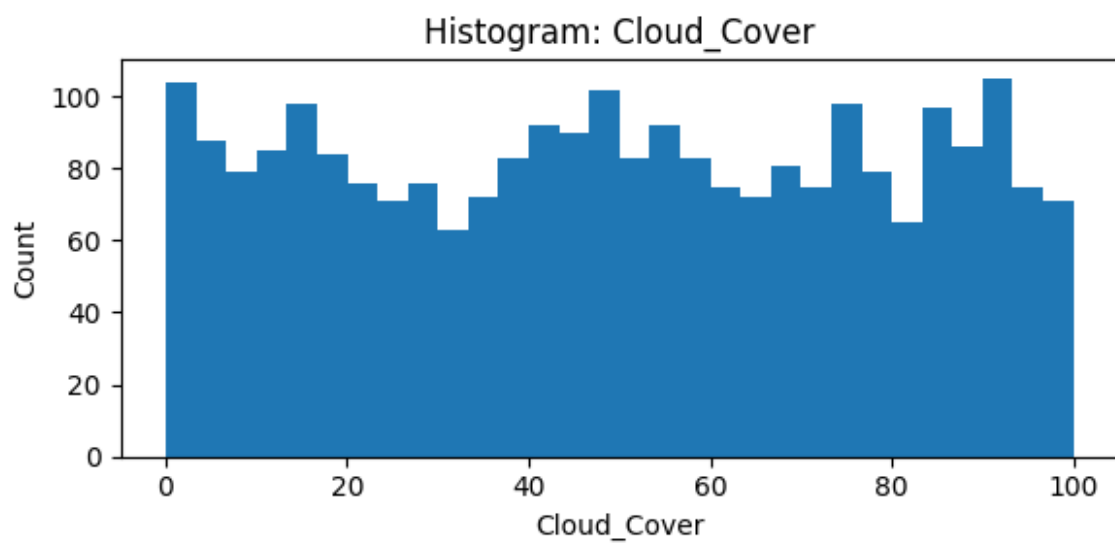
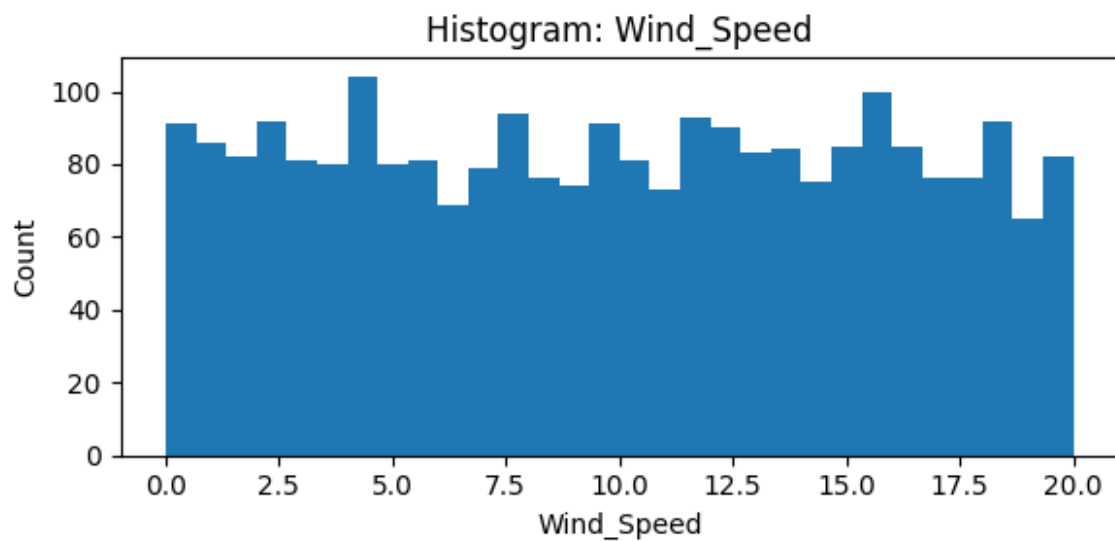
```

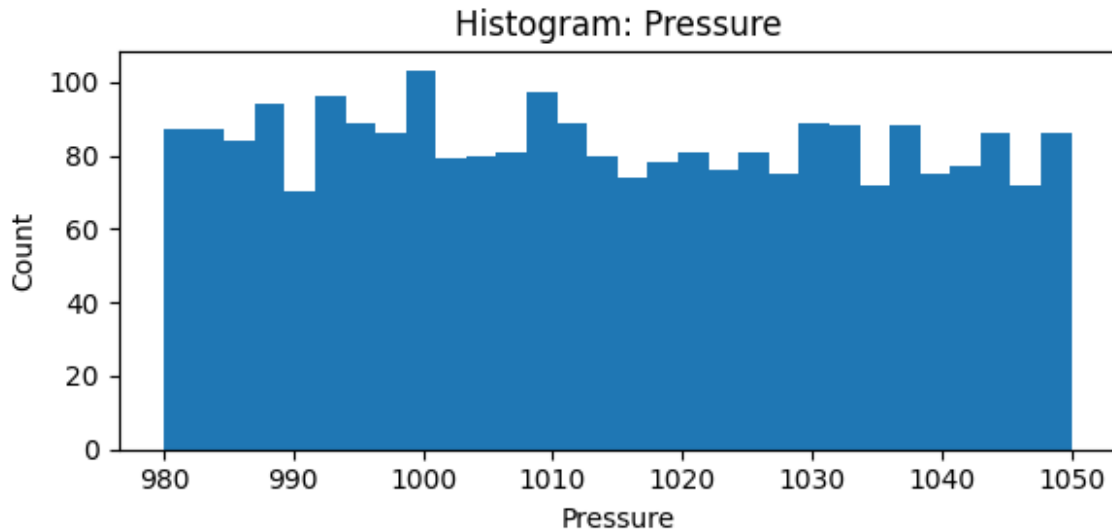
num_cols_non_target = [c for c in num_cols if c != 'Rain_bin']
else:
    num_cols_non_target = num_cols

for col in num_cols_non_target:
    plt.figure(figsize=(6,3))
    plt.hist(df_clean[col].dropna(), bins=30)
    plt.title(f"Histogram: {col}")
    plt.xlabel(col)
    plt.ylabel("Count")
    plt.tight_layout()
    plt.show()

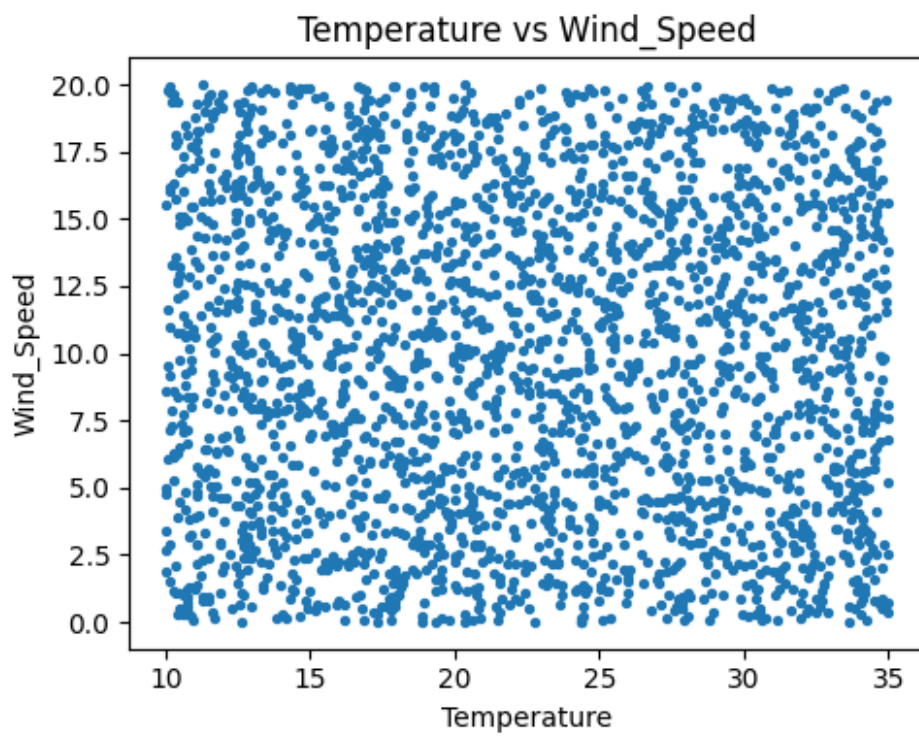

```


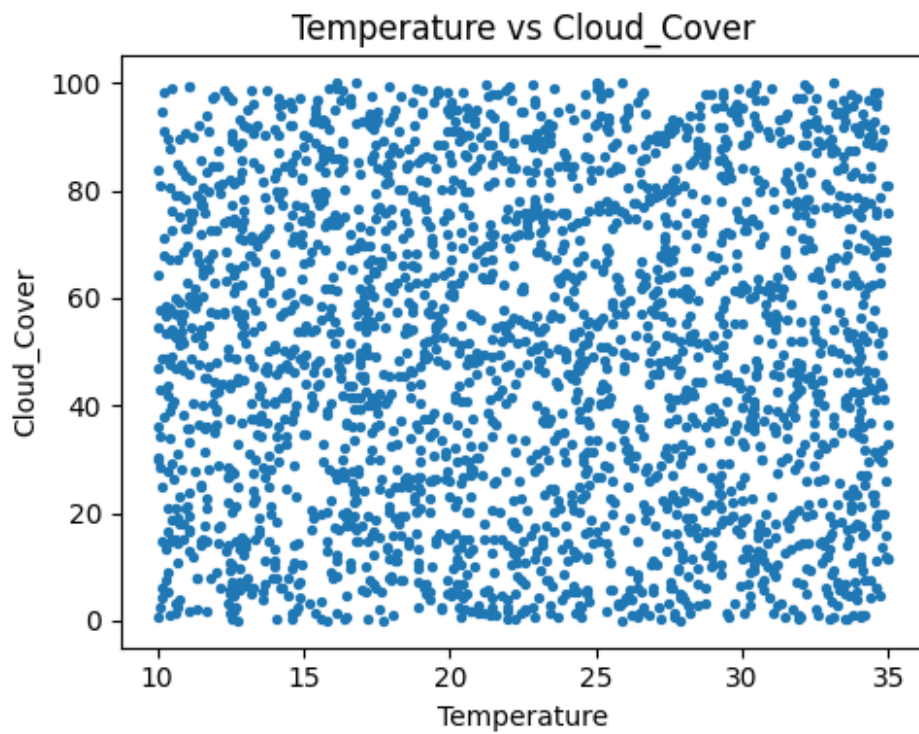


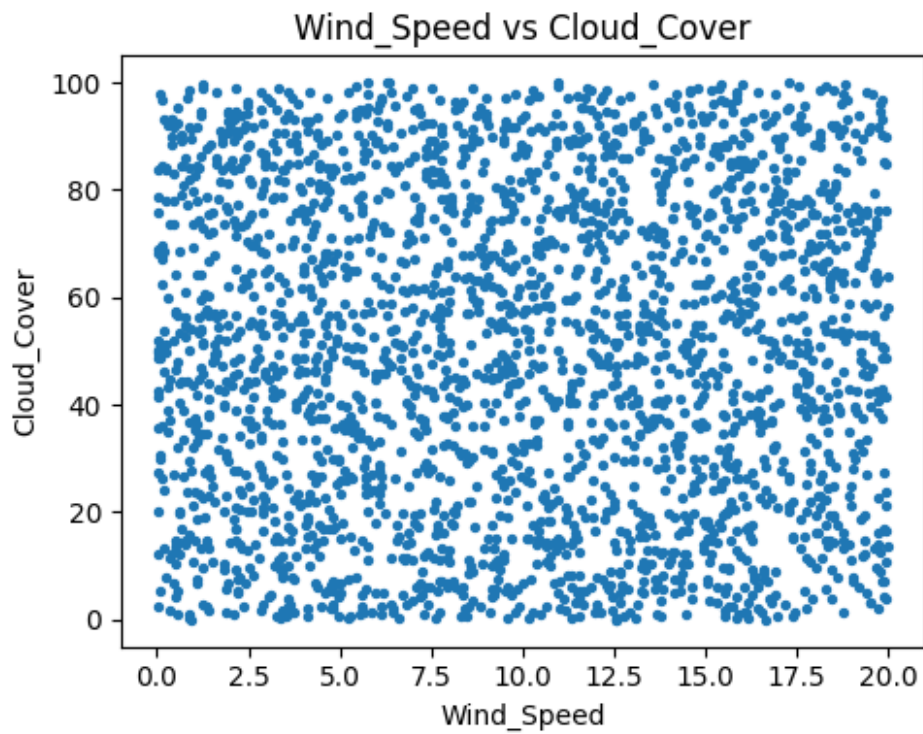
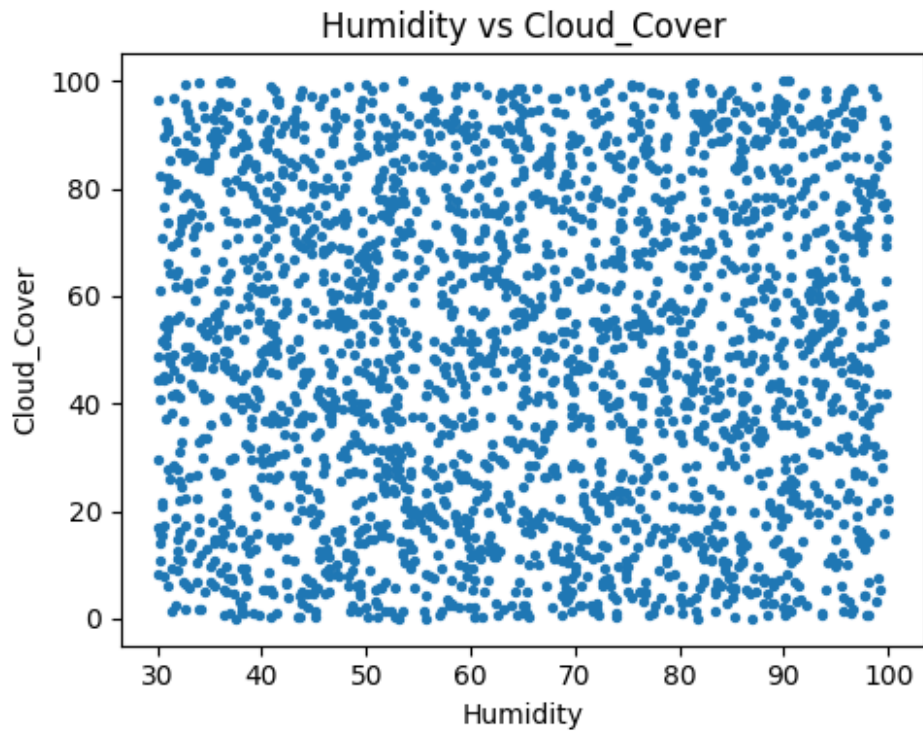




```
# Scatter matrix-ish: pairwise scatter for first 4 numeric features
pairs = num_cols_non_target[:4]
if len(pairs) >= 2:
    for i in range(len(pairs)):
        for j in range(i+1, len(pairs)):
            plt.figure(figsize=(5,4))
            plt.scatter(df_clean[pairs[i]], df_clean[pairs[j]], s=8)
            plt.xlabel(pairs[i])
            plt.ylabel(pairs[j])
            plt.title(f"{pairs[i]} vs {pairs[j]}")
            plt.tight_layout()
            plt.show()
```

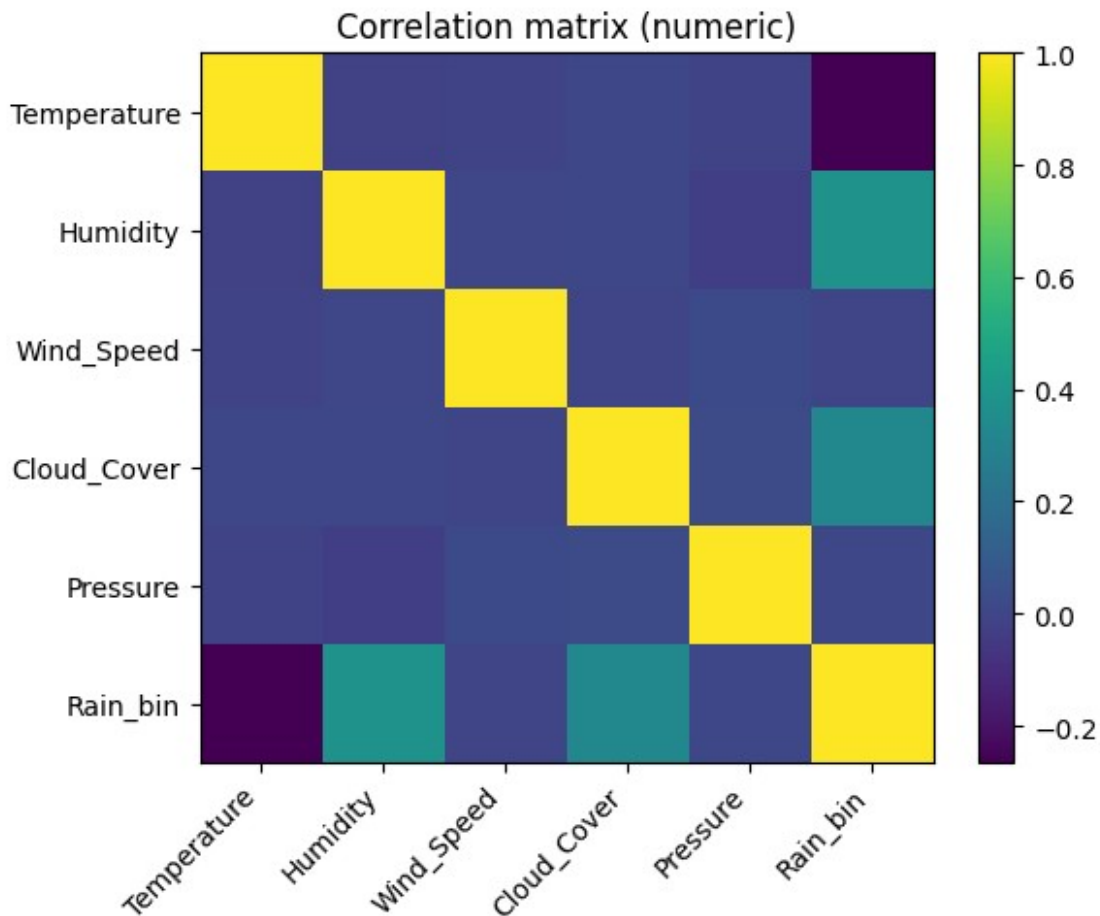






```
# Correlation heatmap using imshow
plt.figure(figsize=(6,5))
plt.imshow(corr, aspect='auto')
plt.colorbar()
```

```
plt.xticks(range(len(corr.columns)), corr.columns, rotation=45,
ha='right')
plt.yticks(range(len(corr.index)), corr.index)
plt.title("Correlation matrix (numeric)")
plt.tight_layout()
plt.show()
```



```
# Modeling: predict Rain_bin
X = df_clean[num_cols_non_target]
y = df_clean['Rain_bin']

# train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)

# scaling
scaler = StandardScaler()
X_train_s = scaler.fit_transform(X_train)
X_test_s = scaler.transform(X_test)
```

```

# Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train_s, y_train)
y_pred = rf.predict(X_test_s)

acc = accuracy_score(y_test, y_pred)
print(f"\nRandom Forest accuracy: {acc:.4f}")
print("\nClassification report:\n", classification_report(y_test,
y_pred, target_names=le.classes_))
print("\nConfusion matrix:\n", confusion_matrix(y_test, y_pred))

# Feature importances
feat_imp = pd.Series(rf.feature_importances_,
index=X.columns).sort_values(ascending=False)
print("\nFeature importances:\n", feat_imp)

# Save model and scaler
joblib.dump(rf, "/mnt/data/rf_weather_model.pkl")
joblib.dump(scaler, "/mnt/data/rf_weather_scaler.pkl")
print("\nSaved model to /mnt/data/rf_weather_model.pkl and scaler
to /mnt/data/rf_weather_scaler.pkl")

# Save a cleaned CSV
df_clean.to_csv("/mnt/data/weather_forecast_data_cleaned.csv",
index=False)
print("Saved cleaned data to
/mnt/data/weather_forecast_data_cleaned.csv")

```

Random Forest accuracy: 1.0000

Classification report:

	precision	recall	f1-score	support
no rain	1.00	1.00	1.00	437
rain	1.00	1.00	1.00	63
accuracy			1.00	500
macro avg	1.00	1.00	1.00	500
weighted avg	1.00	1.00	1.00	500

Confusion matrix:

```
[[437  0]
 [ 0  63]]
```

Feature importances:

```
Humidity      0.376693
Cloud_Cover   0.329689
Temperature   0.265187
Pressure      0.017021
```

```
Wind_Speed      0.011410
dtype: float64
```

```
-----
-----
```

```
FileNotFoundError                                Traceback (most recent call
last)
```

```
Cell In[12], line 28
```

```
    25 print("\nFeature importances:\n", feat_imp)
    27 # Save model and scaler
--> 28 joblib.dump(rf, "/mnt/data/rf_weather_model.pkl")
    29 joblib.dump(scaler, "/mnt/data/rf_weather_scaler.pkl")
    30 print("\nSaved model to /mnt/data/rf_weather_model.pkl and
scaler to /mnt/data/rf_weather_scaler.pkl")
```

```
File ~\AppData\Local\Programs\Python\Python310\lib\site-packages\
joblib\numpy_pickle.py:599, in dump(value, filename, compress,
protocol)
```

```
    597         NumpyPickler(f, protocol=protocol).dump(value)
    598 elif is_filename:
--> 599     with open(filename, "wb") as f:
    600         NumpyPickler(f, protocol=protocol).dump(value)
    601 else:
```

```
FileNotFoundError: [Errno 2] No such file or directory:
'/mnt/data/rf_weather_model.pkl'
```

```
pip install --upgrade --force-reinstall matplotlib scipy pandas
scikit-learn
```

```
pip install --upgrade matplotlib scipy pandas
```

```
pip install --upgrade 'pybind11>=2.12'
```

```
pip install --no-binary :all: matplotlib
```